KP02 Two-Layer Perceptron: Encode Positions First, Then Classify

> Level: undergraduate, year 3
>
> Pattern Recognition | Chapter 4, §4.3

Slides: read and download

Learning objectives

- Explain what the hidden layer and the output layer each do, and turn scores into bits with the step rule.
- Compute the hidden code and the output score of each XOR point and check the predictions.
- Translate the whole two-layer network back to the original plane as a strip between two lines, and explain which merges make the output layer fail.

KP02 Two-Layer Perceptron: Encode Positions First, Then Classify

One-sentence summary: One line cannot answer "A or B?", so first ask two yes-or-no questions with two lines, replace each point by a two-bit code, and then classify with one line of the output layer in the plane of codes. Translated back to the original plane, the whole network is the strip between two parallel lines.

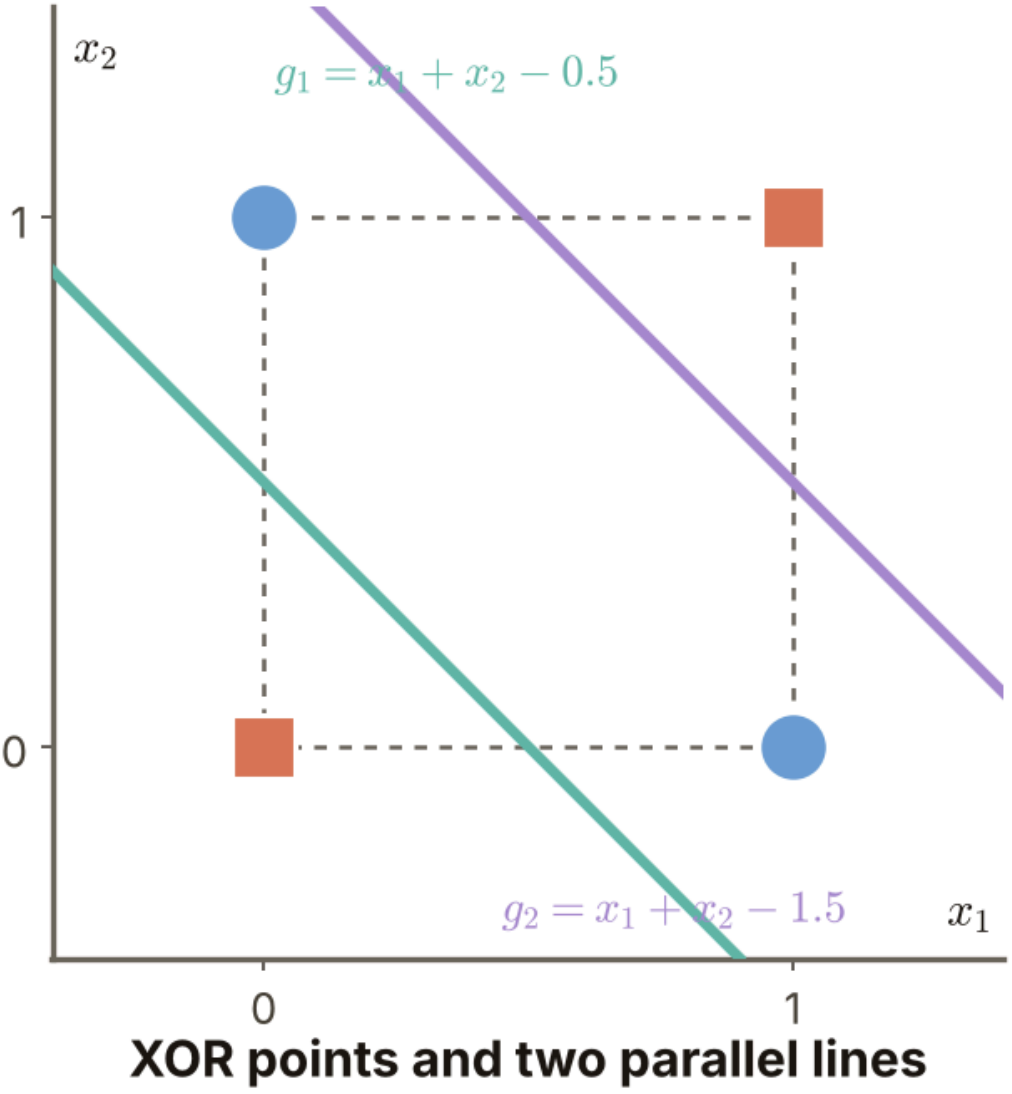
1. One line fails, so ask two small questions first

Plain explanation: The previous KP ended with two parallel lines enclosing a strip that holds exactly the A points of XOR. This section turns that into a network.

- **Hidden layer**: the row of lines that asks the small questions.
- **Output layer**: the last plane that classifies from the codes.

- **Two-layer network**: hidden layer plus output layer; the inputs only supply coordinates and are not a layer.

Figure `lec05-kp02-input-lines`: the four XOR points with the two parallel lines $g_1 = x_1 + x_2 - 0.5$ (cyan) and $g_2 = x_1 + x_2 - 1.5$ (purple).



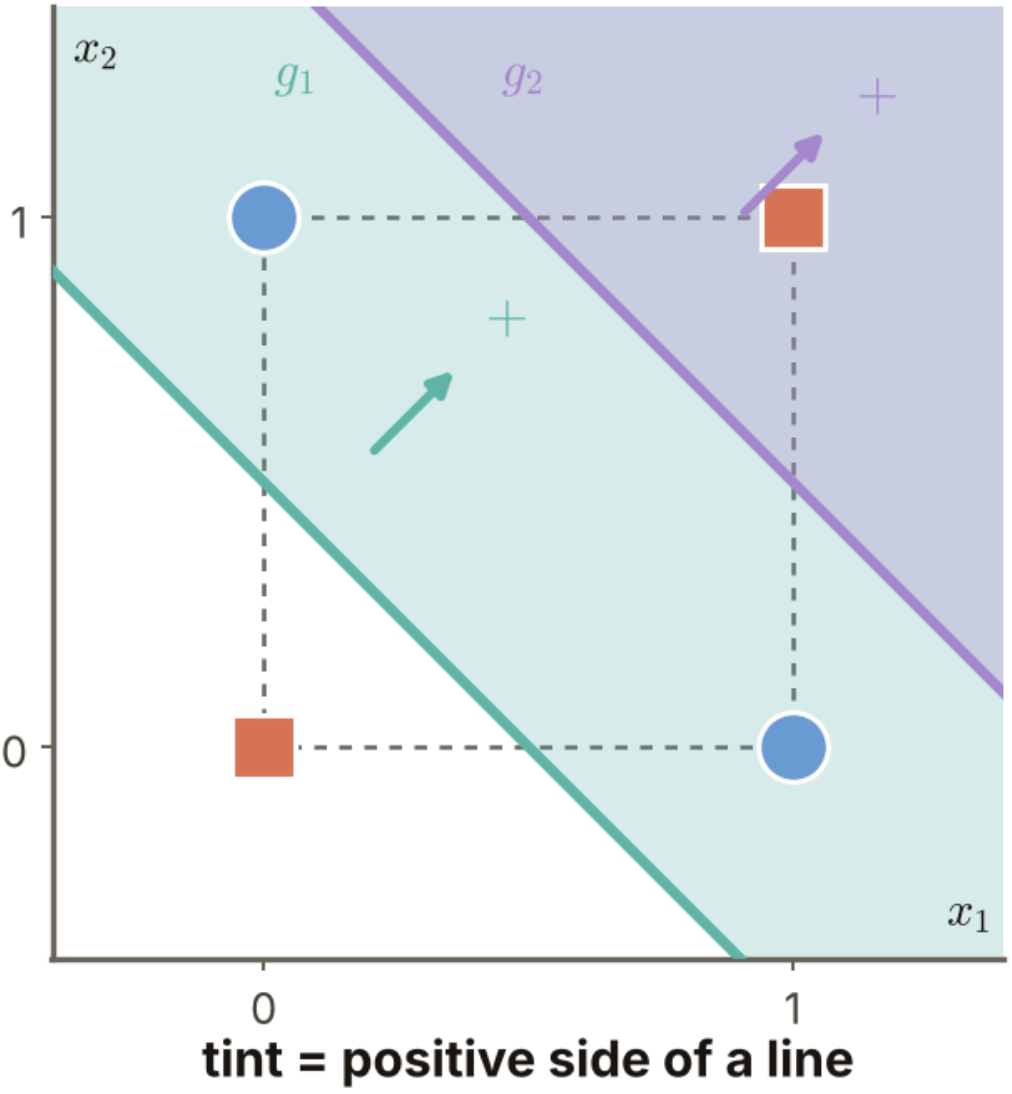
2. Each line asks: is the point on the positive side? Record one bit

Formal definition: for each line g_j, a point on the positive side (score $g_j > 0$) gets 1 and otherwise 0 (exactly 0 also gives 0). This rule is the **step rule**, and the 0 or 1 it produces is a **bit**:

$$y_j = \begin{cases} 1, & g_j(\mathbf{x}) > 0, \\ 0, & g_j(\mathbf{x}) \le 0. \end{cases}$$

The two bits written together, $\mathbf{y} = [y_1, y_2]$, form the **hidden code**, like two position labels attached to the point.

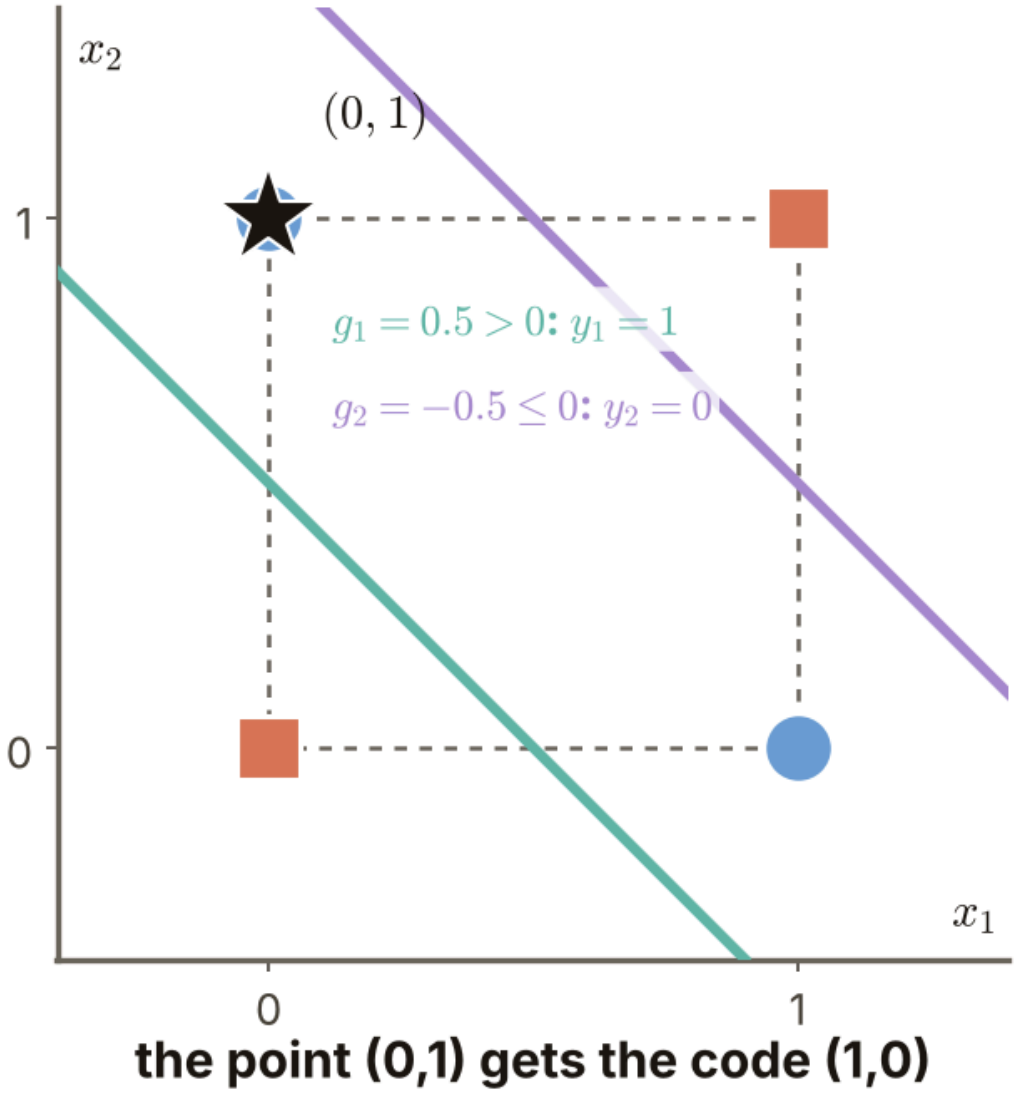
Figure `lec05-kp02-input-sides`: the same picture with the positive side of each line shaded and marked by an arrow; where the colours overlap, both lines are on the positive side.



3. Walk the point $(0, 1)$ through

Example: Into the first line: $g_1 = 0 + 1 - 0.5 = 0.5 > 0$, so $y_1 = 1$. Into the second line: $g_2 = 0 + 1 - 1.5 = -0.5 \leq 0$, so $y_2 = 0$. The hidden code is $(1, 0)$. The flow reads top to bottom: input, two scores, step rule, hidden code.

Figure `lec05-kp02-bits-example`: the star is the point $(0, 1)$, labelled with its two scores and two bits.

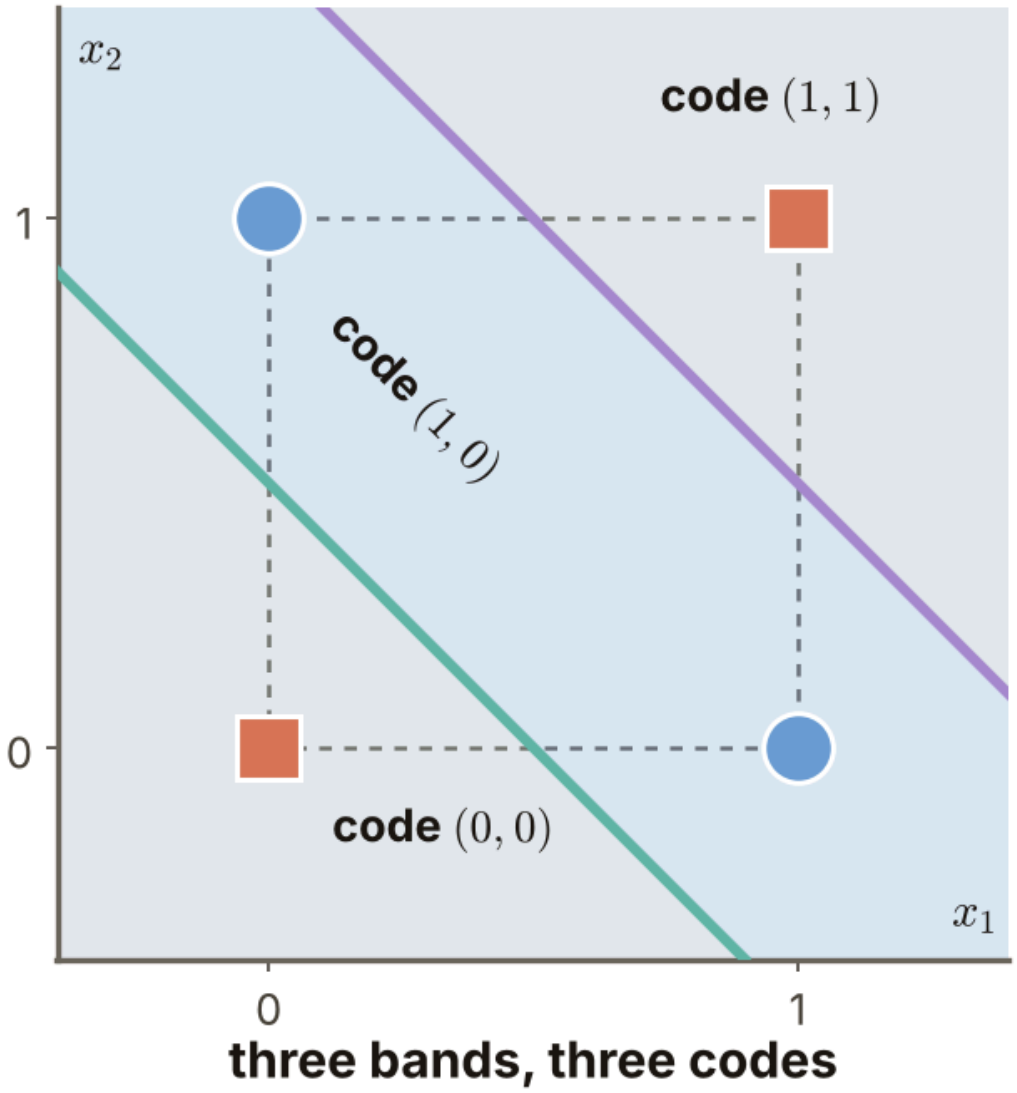


4. The four codes: $(0,0)$, $(1,0)$, $(1,0)$, $(1,1)$

Input	g_1	g_2	Code	XOR truth
$(0,0)$	-0.5	-1.5	$(0,0)$	B
$(0,1)$	0.5	-0.5	$(1,0)$	A
$(1,0)$	0.5	-0.5	$(1,0)$	A
$(1,1)$	1.5	0.5	$(1,1)$	B

The two lines cut the plane into three bands with codes $(0,0)$, $(1,0)$, $(1,1)$; both A points sit in the middle band.

Figure `lec05-kp02-bands`: three bands, each labelled with its code; both A points lie in the $(1,0)$ band.



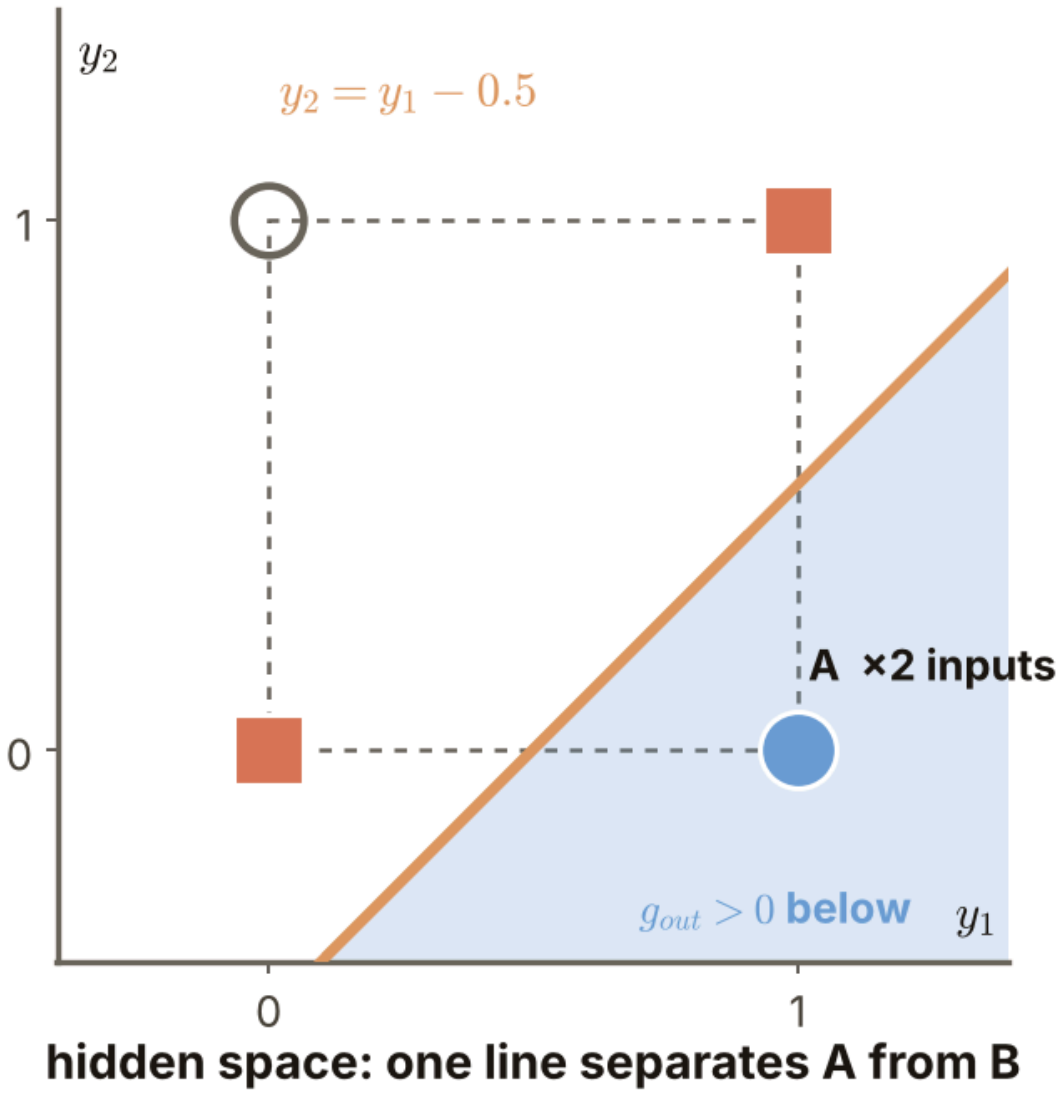
5. In the plane of codes, one line separates A from B

Hidden space: replace every point by its code (y_1, y_2) and draw it in a new plane. $(0, 0)$ is B, $(1, 0)$ is A (two merged inputs) and $(1, 1)$ is B. The output line is

$$g_{\text{out}} = y_1 - y_2 - 0.5, \qquad y_2 = y_1 - 0.5,$$

and $g_{\text{out}} > 0$ (below the line) means A. Merging the two A points does not matter, because they are the same class.

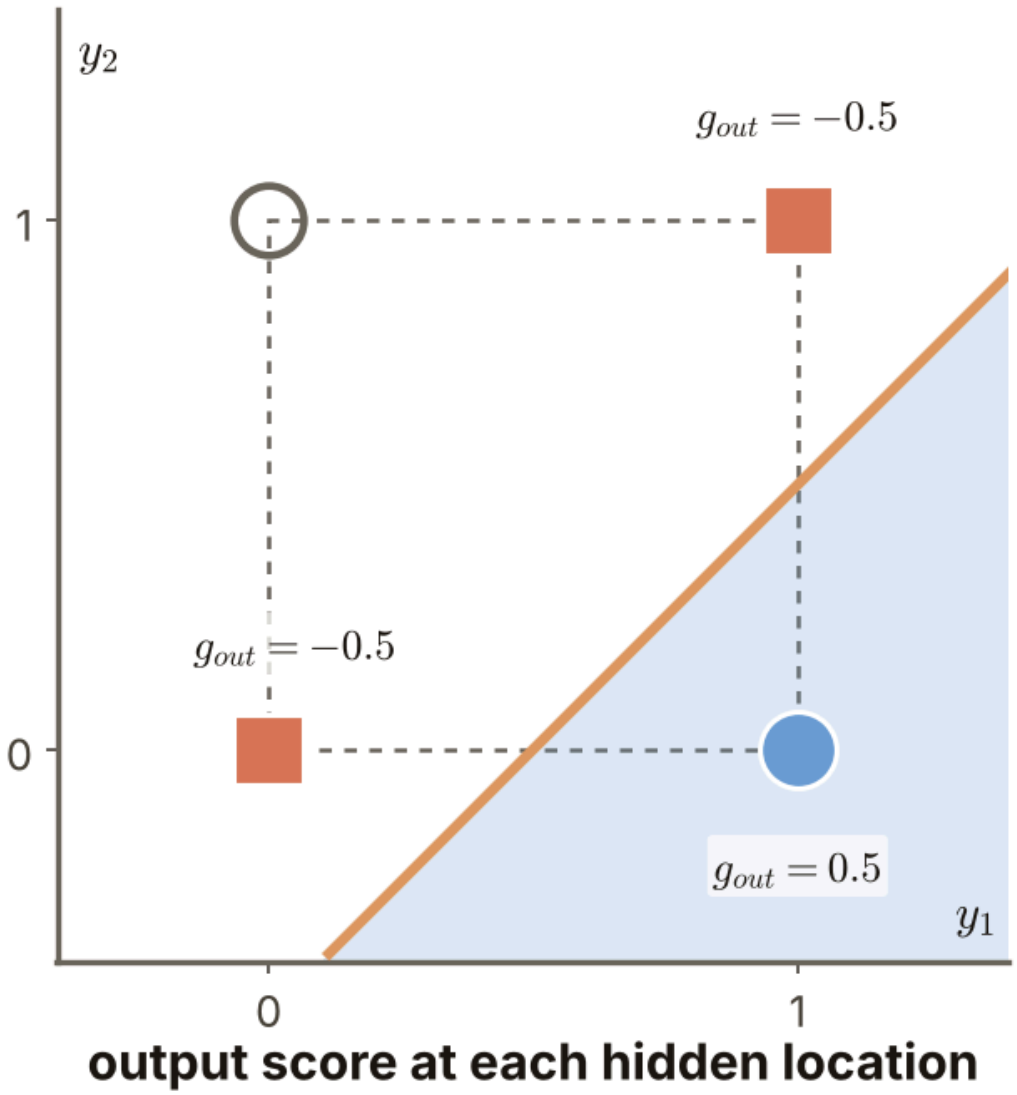
Figure `lec05-kp02-hidden-plane`: the location $(0, 1)$ is hollow because no input lands there; the output line separates $(1, 0)$ from the other two locations.



6. Output layer check: predictions $(0, 1, 1, 0)$ are exactly XOR

$(0, 0)$: $0 - 0 - 0.5 = -0.5$, predict B. $(1, 0)$: $1 - 0 - 0.5 = 0.5$, predict A (both A inputs). $(1, 1)$: $1 - 1 - 0.5 = -0.5$, predict B. The predictions for the four inputs are $(0, 1, 1, 0)$: 4 of 4 correct. This checks these four points; it is not an accuracy on new data.

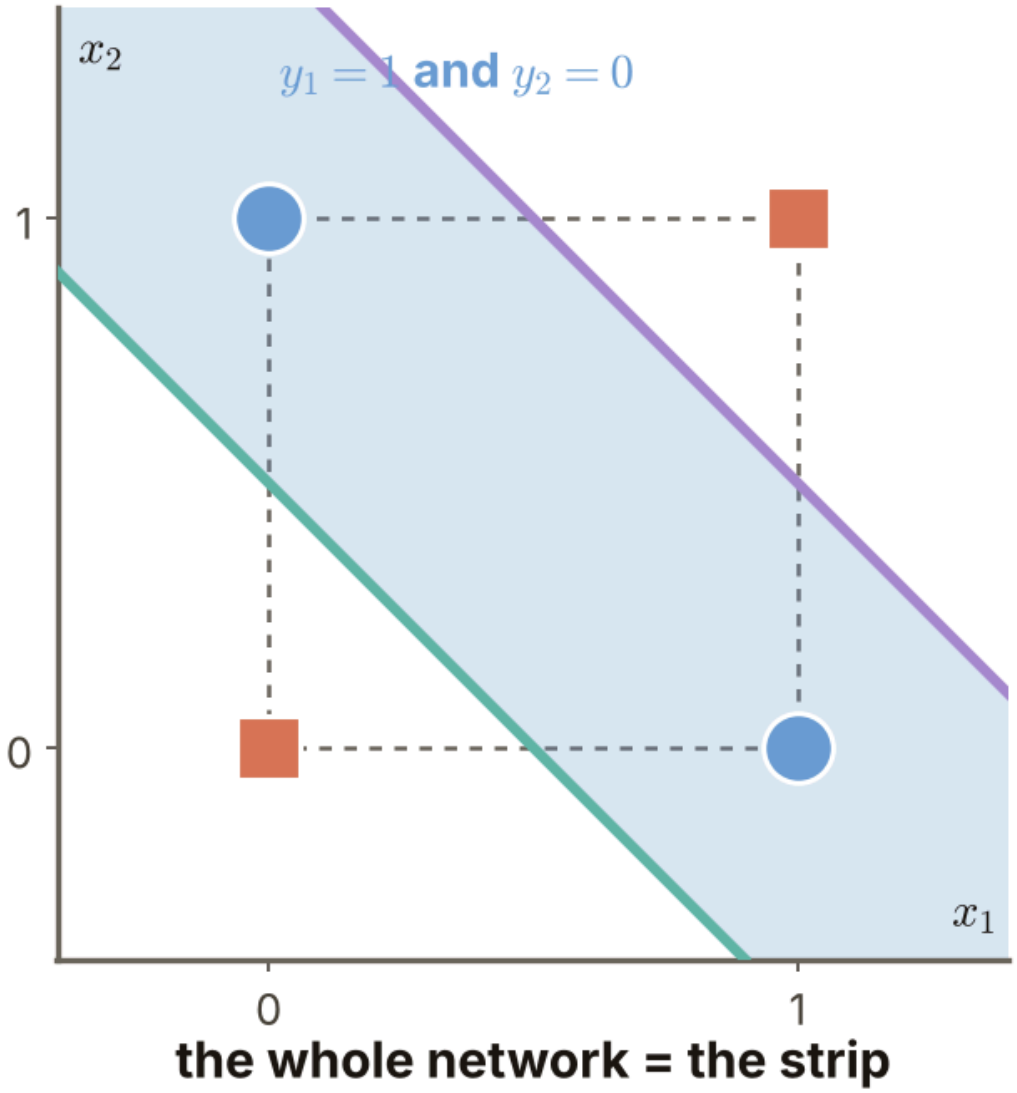
Figure `lec05-kp02-hidden-scores`: the output scores at the three locations; only $(1, 0)$ is positive.



7. Back in the original plane: the whole network is the strip between two lines

$y_1 = 1$ and $y_2 = 0$ means $x_1 + x_2 > 0.5$ and $x_1 + x_2 \leq 1.5$. This is the strip of the previous KP: the hidden layer makes the two decisions and the output layer asks for both.

Figure `lec05-kp02-strip-net`: both A points inside the strip, both B points outside.

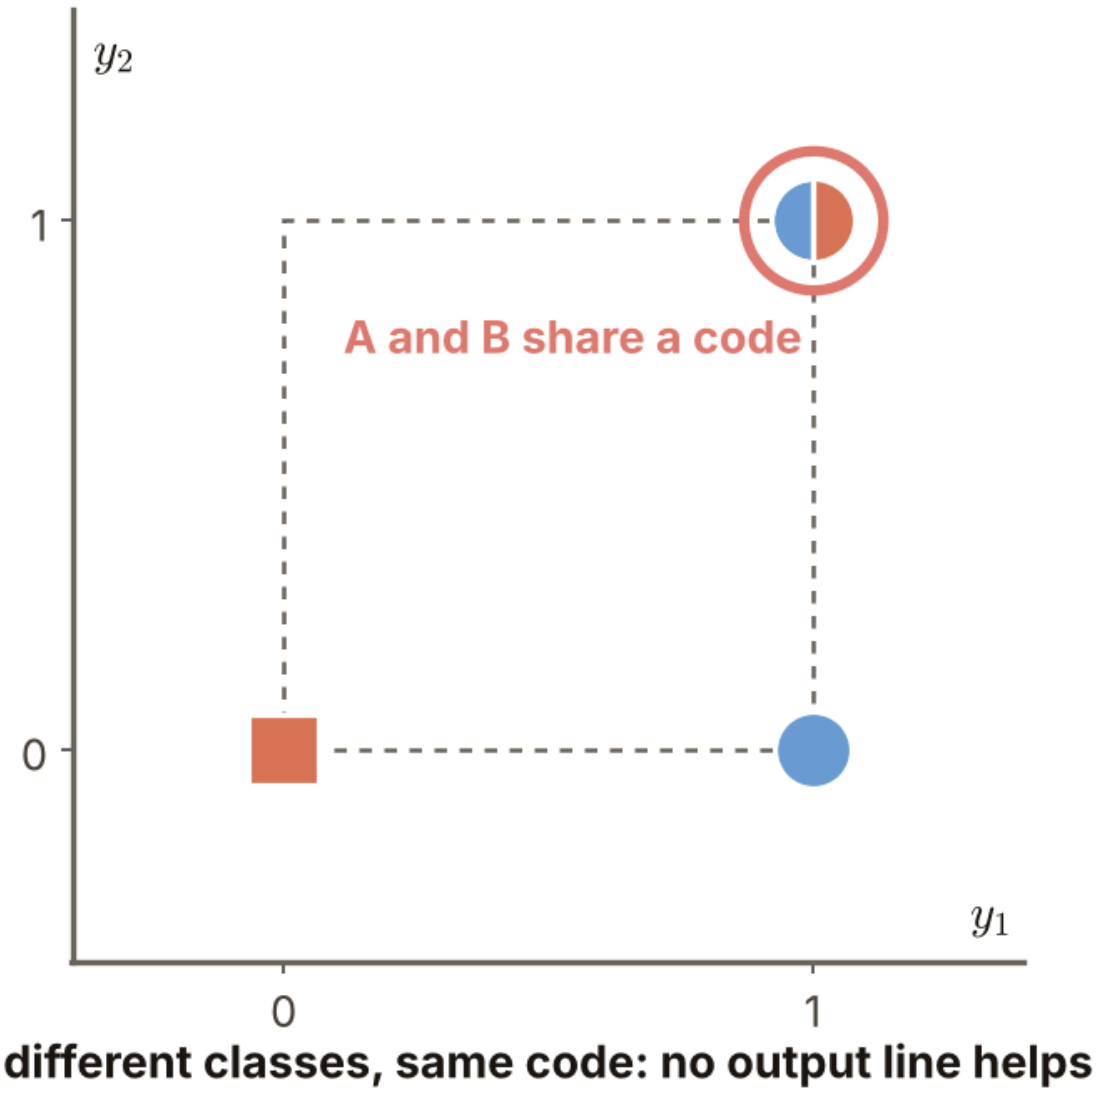


8. Merging the same class is fine; merging different classes is not

Here the two A points merge and the output layer still separates them.

Common misconception: treating the hidden bits as class probabilities, or assuming any merge is harmless. But if points of different classes receive the same code, no output line can separate them and later decisions cannot recover. Note that y_1, y_2 are only "on the positive side?" bits, not class probabilities. This construction only proves that XOR can be represented; for other tasks, first check that the hidden codes can be separated by the output layer.

Figure `lec05-kp02-merge-bad`: schematic with an A and a B both at $(1, 1)$, the conflict marked by a red ring.



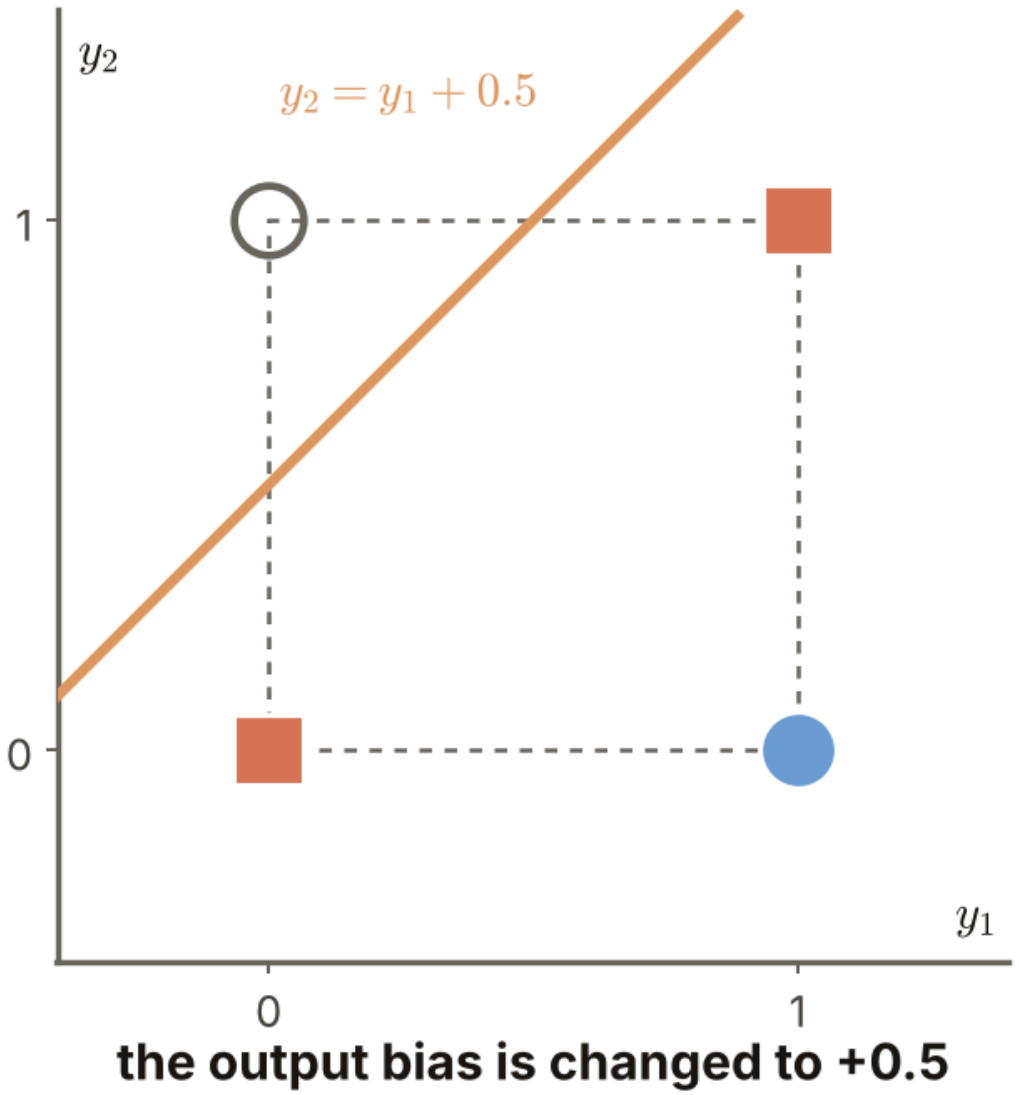
The next KP asks: with more hidden lines, can any labelling be separated?

Extension: practice and worked solution

Practice: keep the hidden codes and change the output score to $g_{\text{out}} = y_1 - y_2 + 0.5$.

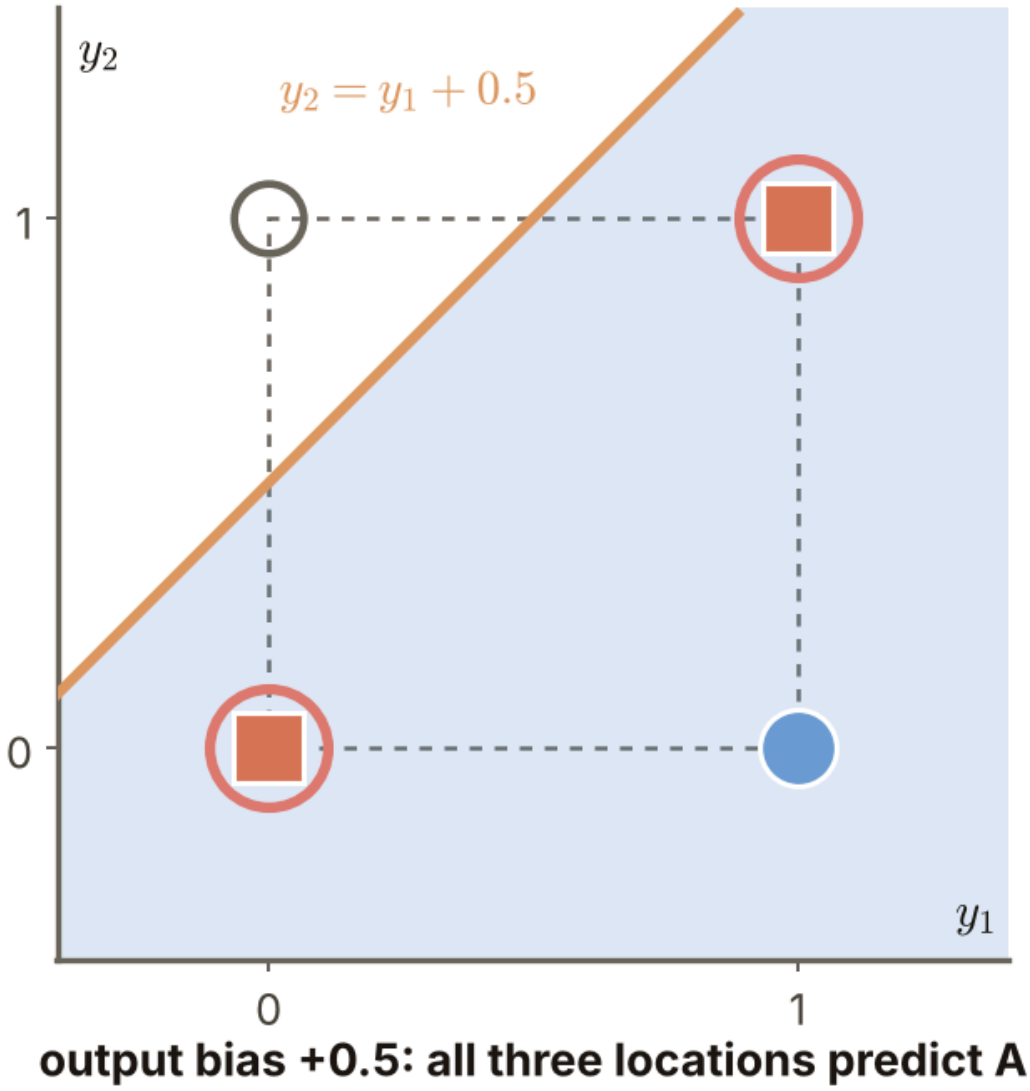
1. What are the predictions for the four inputs? Hint: recompute only the output scores.
2. Which different inputs are merged, and why does the merge cause no conflict? Hint: look at their labels.

Figure `lec05-kp02-bias-line`: the three locations and the new output line $y_2 = y_1 + 0.5$.



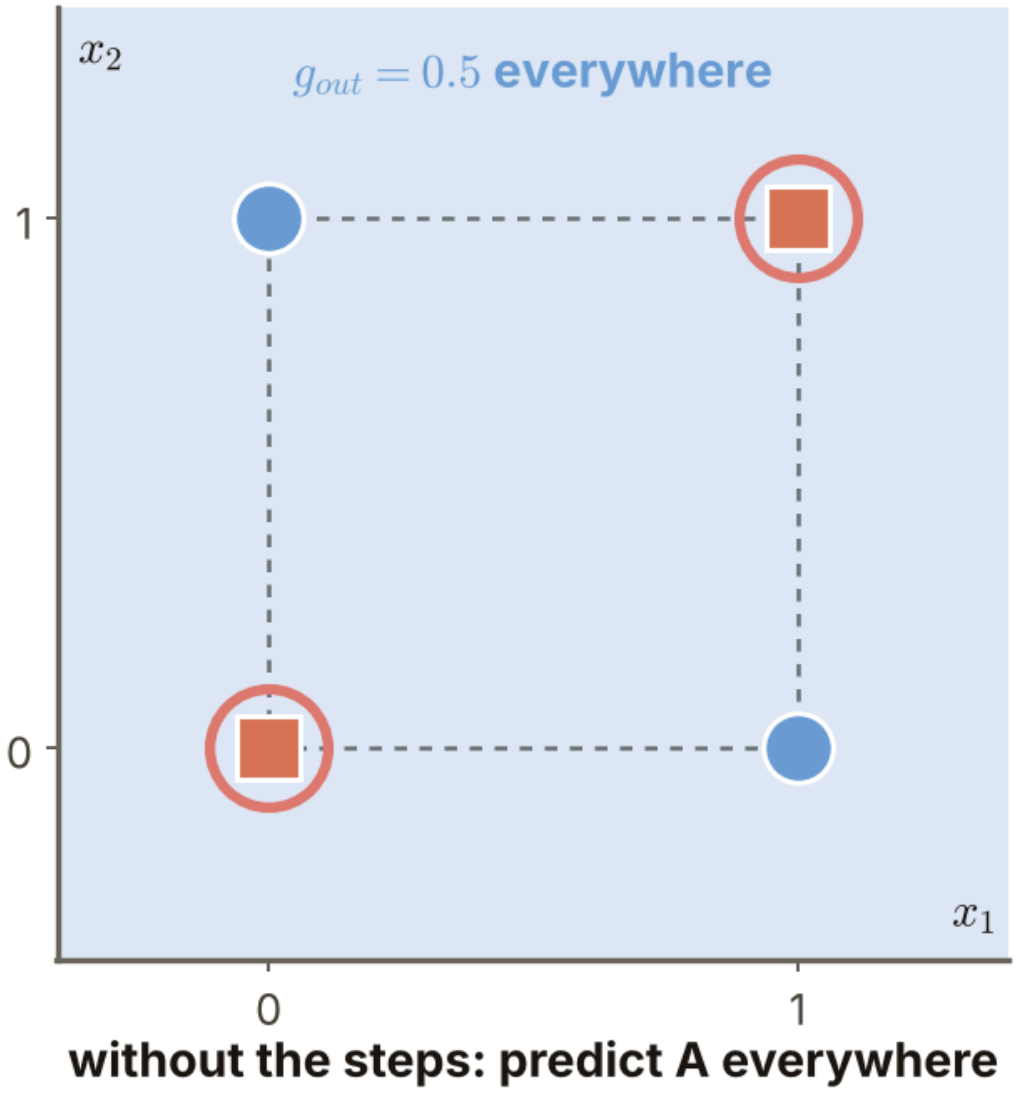
Worked solution of question 1: (1) The formula is $g_{\text{out}} = y_1 - y_2 + 0.5$. (2) Substitute the three locations: $(0, 0)$: 0.5; $(1, 0)$: 1.5; $(1, 1)$: 0.5. (3) All three scores are positive, so all four inputs are predicted A. (4) Check: XOR is B, A, A, B, so the two B inputs are predicted A: 2 errors. The bias shifted the line so it no longer separates B from A.

Figure `lec05-kp02-bias-plus`: the two B locations marked by red rings as errors.



The steps cannot be removed [added]: without the steps, $y_1 = x_1 + x_2 - 0.5$ and $y_2 = x_1 + x_2 - 1.5$; substituting into $g_{\text{out}} = y_1 - y_2 - 0.5$ gives $(x_1 + x_2 - 0.5) - (x_1 + x_2 - 1.5) - 0.5 = 0.5$, independent of x, so every point is predicted A and XOR has 2 errors. A combination of linear expressions is linear again, so the nonlinearity has to come from the steps.

Figure `lec05-kp02-no-step`: the whole plane predicts A, the two B points marked by red rings.



Terms and sources

English	中文
Hidden layer, output layer	隱藏層、輸出層
Bit	位元
Step rule	階梯規則
Hidden code	隱藏編碼
Hidden space	隱藏空間

Source: Chapter 4 KP02, §4.3 (read 2026-10-24). The two-layer XOR construction follows the source; explanations and figures were rewritten; the bias practice and the no-step derivation are [added].